

Wisdom Science API Longitudinal Panel: Six-Model Low-Cost Evidence Layer

Mian Zhang

Abstract

The API longitudinal panel is a low-cost, local evidence layer for Wisdom Science. It evaluates multiple text models over repeated WisdomBench rounds, records raw responses, scores them with a fixed LLM judge, and reports first-round intelligence, final-round performance, WQ, and repeat failure rate. It does not replace embodied rollout evidence. Its purpose is to test whether the intelligence-wisdom separation, memory mechanisms, and evidence gates survive across a broader model panel before expensive cloud or robot experiments are launched.

1 Protocol

The panel uses six API model endpoints, four longitudinal strategies, WisdomBench tasks, fixed seeds, and repeated rounds. Each cell records the model endpoint, strategy, task, seed, round, prompt hash, response hash, raw response, judge model, score, and judge reason. API credentials are never written to disk; the config stores only environment variable names.

- **Models:** six OpenAI-compatible endpoints from Claude, DeepSeek, and Qwen families when credentials are available.
- **Strategies:** no memory, self-refine, Reflexion, and Cognitive Immunity.
- **Metrics:** first-round score, final-round score, WQ, RFR, raw row count, and evidence-gate status.
- **Boundary:** API text evidence cannot be reported as robot rollout or VLA policy evidence.

2 Evidence Gate

A panel row is admissible only when it includes a config hash, model endpoint identity, judge identity, task ID, seed, round, prompt hash, response hash, raw response, score, and judge reason. A dry run is useful for readiness validation but is not empirical evidence. Rows with API errors remain in the raw log and are not silently dropped.

3 Current Table

The table below is generated only when the runner has produced raw rows. If it is empty or dry-run only, the manuscript must describe the panel as prepared rather than completed.

Table 1: API WisdomBench longitudinal panel summary.

Model	Strategy	Tasks	R1	Rfinal	Delta	WQ
Claude Opus Proxy	cognitive_immunity	2	3.00	2.00	-1.00	0.000
Claude Opus Proxy	no_memory	2	3.00	3.00	+0.00	0.000
DeepSeek Chat	cognitive_immunity	2	2.50	3.00	+0.50	0.500
DeepSeek Chat	no_memory	2	2.50	2.50	+0.00	0.000
DeepSeek Reasoner	cognitive_immunity	2	2.50	3.00	+0.50	0.500
DeepSeek Reasoner	no_memory	2	2.00	2.50	+0.50	0.250
Qwen Max	cognitive_immunity	2	1.00	3.00	+2.00	1.000
Qwen Max	no_memory	2	1.00	0.50	-0.50	-0.500
Qwen Plus	cognitive_immunity	2	2.50	3.00	+0.50	0.500
Qwen Plus	no_memory	2	2.50	2.50	+0.00	0.000
Qwen Turbo	cognitive_immunity	2	2.50	1.50	-1.00	0.500
Qwen Turbo	no_memory	2	1.00	1.00	+0.00	0.000

4 Use in the Portfolio

This panel is the cheapest next empirical layer above the current formal, physics, and representation package. It strengthens the claim that Wisdom Science is not tied to one model family, while preserving the evidence boundary: text API panels support longitudinal cognitive claims, not embodied rollout claims.